


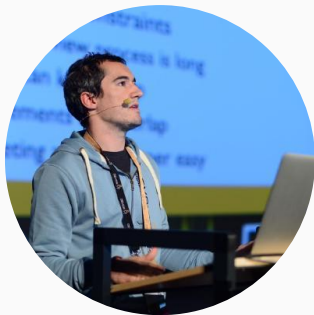
Observability: buzz-word or real need?

<https://github.com/montrealjug/jug-montreal-otel>

Maxime DAVID & Olivier GATIMEL
Montreal JUG - Jan 2023



whoarewe >



Maxime David
Serverless @Datadog

AWS Community Builder
Youtuber (rust)



Olivier Gatimel
Lead dev @CARL Berger-Levrault



Agenda

1. Overview
2. Definitions, concepts and use cases
3. Demo
4. Going further
5. Questions!

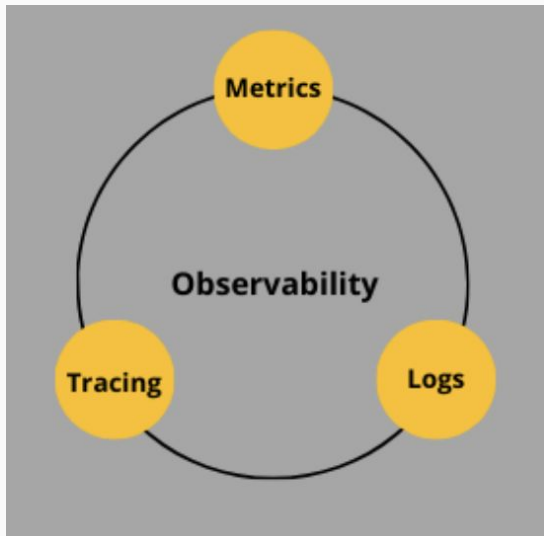
Observability != Monitoring

Observability refers to the ability to understand the internal state of a system by examining its outputs, such as **logs, metrics and traces**. It allows for the diagnosis of issues by providing insight into the system's behavior over time.

Monitoring refers to the continuous collection of data from a system to **check for any abnormal behavior or performance issues**

Without observability -> no monitoring!

Three pillars of observability



- **Logs** provide a record of events that occur within a system
- **Metrics** provide measurable values that can be used to track the performance and health of a system.
- **Traces** provide a detailed record of the steps taken by a request or process as it flows through a distributed system, and can be used for debugging and performance analysis.

Source: <https://iamondemand.com/blog/the-3-pillars-of-system-observability-logs-metrics-and-tracing/>

Logs - Example

Plain text (apache access log)

```
10.1.2.3 - rehg [20/Jan/2023:19:22:12 -0000] "GET /hiMontrealJug HTTP/1.1" 200 3423
```

Logs - Example

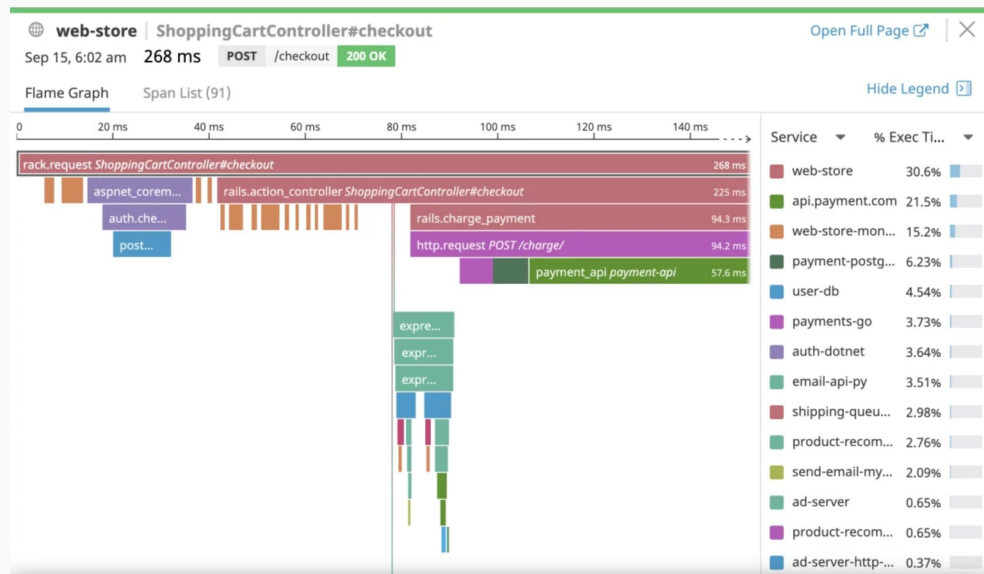
JSON (GCP example)

```
{
  "insertId": "42",
  "jsonPayload": {
    "message": "There was an error in the application",
    "times": "2019-10-12T07:20:50.52Z"
  },
  "httpRequest": {
    "requestMethod": "GET"
  },
  "resource": {
    "type": "k8s_container",
    "labels": {
      "container_name": "hello-app",
      "pod_name": "helloworld-gke-6cfd6f4599-9wff8",
      "project_id": "stackdriver-sandbox-92334288",
      "namespace_name": "default",
      "location": "us-west4",
      "cluster_name": "helloworld-gke"
    }
  },
  "timestamp": "2020-11-07T15:57:35.945508391Z",
  "severity": "ERROR",
  "labels": {
    "user_label_2": "value_2",
    "user_label_1": "value_1"
  },
  "logName": "projects/stackdriver-sandbox-92334288/logs/stdout",
  "operation": {
    "id": "get_data",
    "producer": "github.com/MyProject/MyApplication",
    "first": true
  },
  "sourceLocation": {
    "file": "get_data.py",
    "line": "142",
    "function": "getData"
  },
  "receiveTimestamp": "2020-11-07T15:57:42.411414059Z"
}
```

Metrics - Example

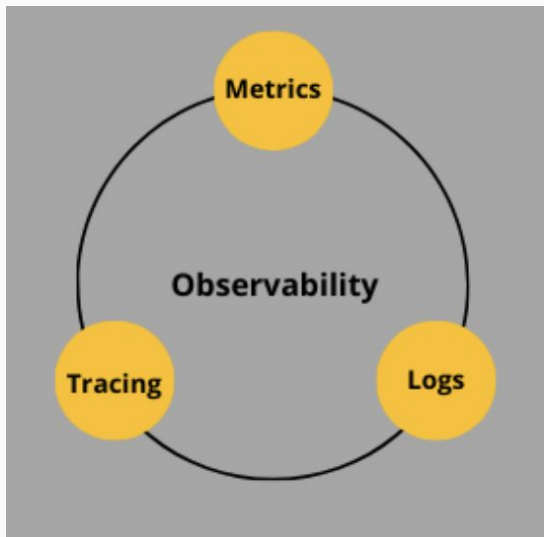
- System metrics
 - `system.cpu.idle`
 - `redis.keys.evicted`
- Business metrics
 - `my.project.cart.item.added`
 - `My.project.password.forget.co`
- OpenTelemetry metric data models
 - Sum (or counter), Gauge, Histogram
 - Exemplar : associates metric with a trace or a span

Traces - Example



Source: <https://docs.datadoghq.com/fr/video-categories/flamegraph/>

Three pillars of observability



Source: <https://iamondemand.com/blog/the-3-pillars-of-system-observability-logs-metrics-and-tracing/>

OpenTelemetry

OpenTelemetry is a collection of tools, APIs, and SDKs.

Use it to **instrument, generate, collect, and export** telemetry data (metrics, logs, and traces) to help you analyze your software's performance and behavior.

Source: <https://opentelemetry.io/>

NOT to store **NOR** visualize data

OpenTelemetry history

- Started in mid 2019 as a sandbox merge project of OpenTracing (traces) and OpenCensus (metrics)
- Supported by Cloud Native Computing Foundation (<https://www.cncf.io/>)
- Switched to incubating project in mid 2021
- Second most important project in CNCF (after Kubernetes)

Is it ready?

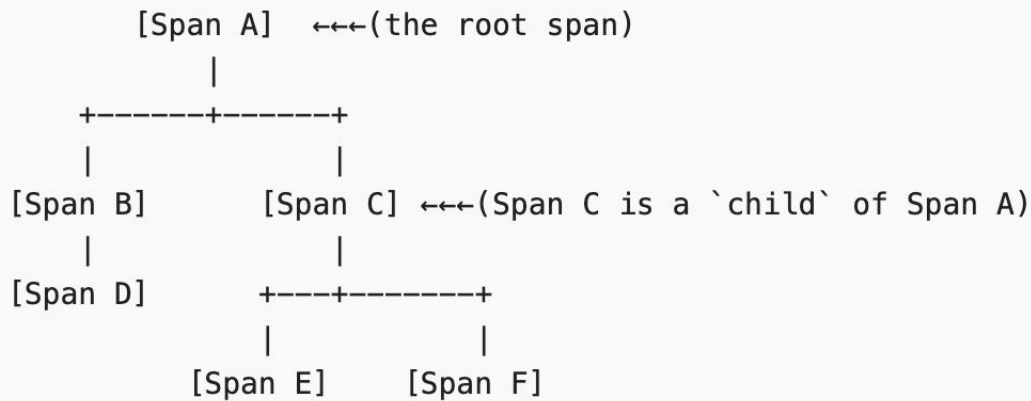
Current status : <https://opentelemetry.io/status> (moving fast)

- **Tracing**
 - Stable with LTS. Version 1.x out since Feb 2021
- **Metric**
 - API and protocol stable since Nov 2021
 - SDK still under active dev, but Meter is ready in java-sdk
- **Logs**
 - Draft/experimental
 - Current focus is integration with existing framework
 - An API will come later

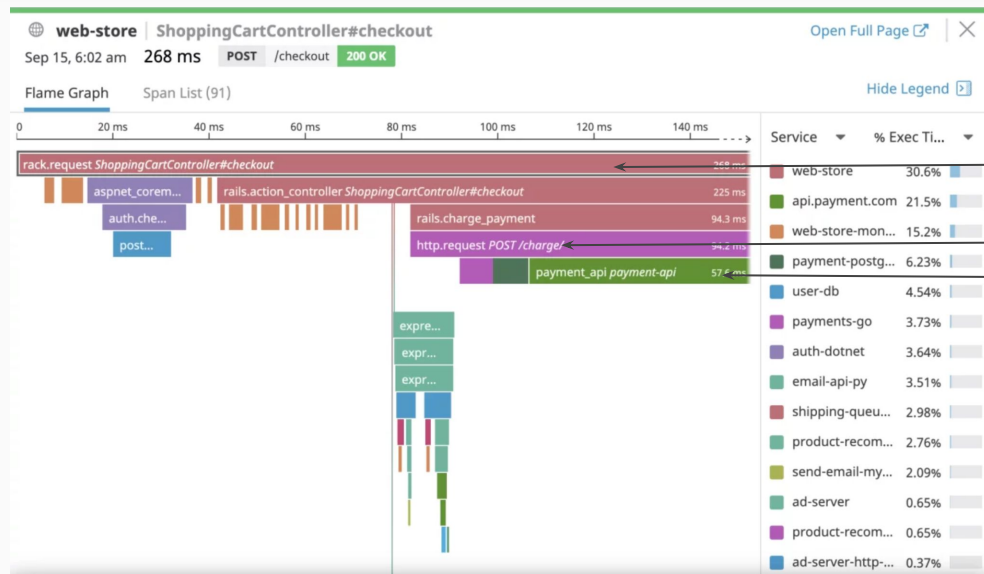
Java is usually the language to get things implemented first

Focus on Tracing -> (more definitions 🙈)

Traces can be viewed as a directed acyclic graph of **Spans**



Focus on Tracing



root span

(child) spans

Source: <https://docs.datadoghq.com/fr/video-categories/flamegraph/>

Spans

Span = operation within a transaction.

Span contains

- Parent's Span identifier (remember the DAG)
- An operation name
- A start and finish timestamp
- Attributes -> key-value pairs.
- A set of zero or more Events, each of which is itself a tuple (timestamp, name, attributes)

```
Span #2
Trace ID      : 75e9a2a6eb8482613901c261d8cf6428
Parent ID     : 98e14f3fbcd9be2
ID            : 1f19681c573efdf5
Name          : HTTP POST
Kind          : Client
Start time    : 2023-01-23 20:05:40.243723707 +0000 UTC
End time      : 2023-01-23 20:05:40.245862987 +0000 UTC
Status code   : Error
Attributes:
-> http.method: Str(POST)
-> http.url: Str(http://localhost:8080/person/carolyn)
Events:
SpanEvent #0
-> Name: exception
-> Timestamp: 2023-01-23 20:05:40.245746748 +0000 UTC
-> Attributes:
-> exception.stacktrace: Str(java.net.ConnectException)
```


Spans (going further)

- *Links to zero or more causally-related Spans (via the SpanContext of those related Spans).*
- *SpanContext information required to reference a Span*

Context propagation

By default in java-instrumentation,

W3C Trace Context (<https://www.w3.org/TR/trace-context/>)

-> published in november 2021

```
traceparent: 00-0af7651916cd43dd8448eb211c80319c-b7ad6b7169203331-01
```

version - traceId - parentId - traceFlags

Alternatives are possible (jaeger, xray, ...)

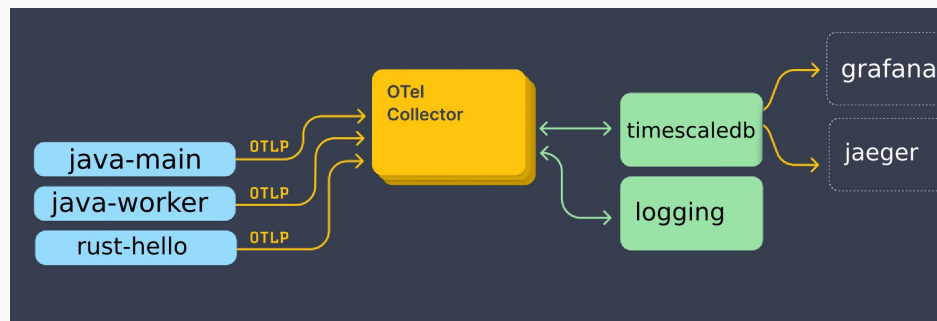
Demo deployment

Receivers:

- java-main: plain Java with sdk
- java-worker: SpringBoot with api + agent
- rust-hello : with sdk

Exporters:

- logging : otel collector stdout exporter
- timescaledb : postgres extension for time series storage
- grafana + jaeger : for visualization



DEMO 🔥

Going further

More than 3 pillars

- Profiling
- Code integration

Baggage

- A baggage is an attribute which is propagated between spans (e.g. article code)
- It is a public information, so caution with its visibility
- W3C Baggage (<https://www.w3.org/TR/baggage/>)
-> working draft since september 2022

Visualization

- You likely need a tool to visualize

Links

Migration

- From OpenTracing
<https://github.com/open-telemetry/opentelemetry-java/blob/main/opentracing-shim>
- From OpenCensus
<https://github.com/open-telemetry/opentelemetry-java/blob/main/opencensus-shim>
- From micrometer with Java agent
<https://github.com/open-telemetry/opentelemetry-java-instrumentation/tree/main/instrumentation/micrometer/micrometer-1.5>
- Send OTEL metrics with micrometer
<https://github.com/open-telemetry/opentelemetry-java-contrib/tree/main/micrometer-meter-provider>
- Use Micrometer to ship metrics in OTLP format
<https://github.com/micrometer-metrics/micrometer/issues/2864>

Links

Java agent alternative ?

Add java-sdk dependency and initialize AutoConfigure

<https://github.com/open-telemetry/opentelemetry-java-docs/blob/main/autoconfigure/src/main/java/io/opentelemetry/example/autoconfigure/AutoConfigExample.java>

Show trace+span id into logging with MDC

<https://github.com/open-telemetry/opentelemetry-java-instrumentation/blob/main/docs/logger-mdc-instrumentation.md>

Testing with traces

<https://github.com/kubeshop/tracetest>

QUESTIONS ?

